

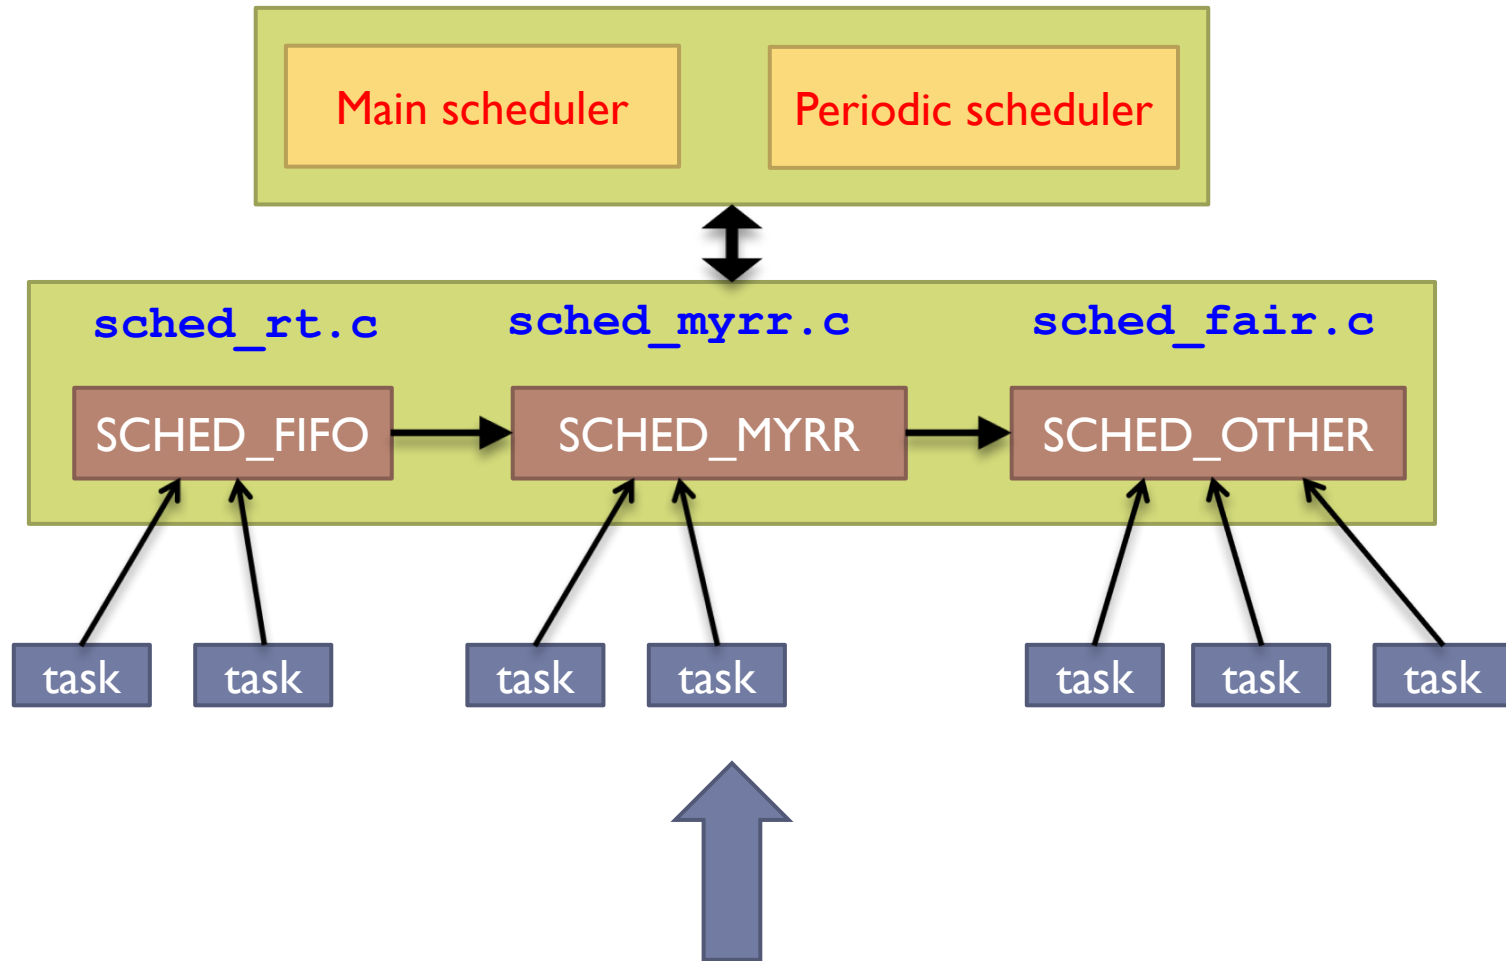
Implement a Round Robin Scheduler

# What is Round-robin scheduling

---

- ▶ Round-robin scheduler gives each job a time slice
  - ▶ When we implement the entity class (i.e. `sched_myrr_entity`), a new attribute called `time_slice` is added
  - ▶ When
    - ▶ The given `time_slice` is expired **AND**
    - ▶ There exists some other tasks in the queue
    - ▶ The current task will be moved to the end of the queue

# SCHED\_MYRR



Now, we replaced MYFIFO with MYRR

# Implementation of SCHED\_MYRR

# Step-1: define scheduling policy no.

---

- ▶ Locate the following line

```
#define SCHED_IDLE 5
```

in `include/uapi/linux/sched.h`

- ▶ After the above line, add the following line

```
#define SCHED_MYRR 6
```

- ▶ If you want to keep the previous MY\_FIFO scheduler policy number, you can set SCHED\_MYRR as 7

## Step-2: define the scheduler entity

---

- ▶ In `include/linux/sched.h`, locate the declaration of `struct sched_rt_entity`, after which we define `struct sched_myrr_entity` to encapsulate needed information to help implement the new scheduling policy

```
struct sched_myrr_entity {  
    struct list_head run_list;  
    int time_slice; /* new field is introduced */  
};
```

- ▶ Locate the definition of `struct task_struct`, add the following line

```
struct sched_myrr_entity myrr;
```

after the line

```
struct sched_rt_entity rt;
```

## Step-3: define a new runqueue

---

- ▶ In `kernel/sched/sched.h`, we define `struct myrr_rq`

```
struct myrr_rq {  
    struct list_head queue;  
};
```

- ▶ Locate the definition of `struct rq`, add the following line

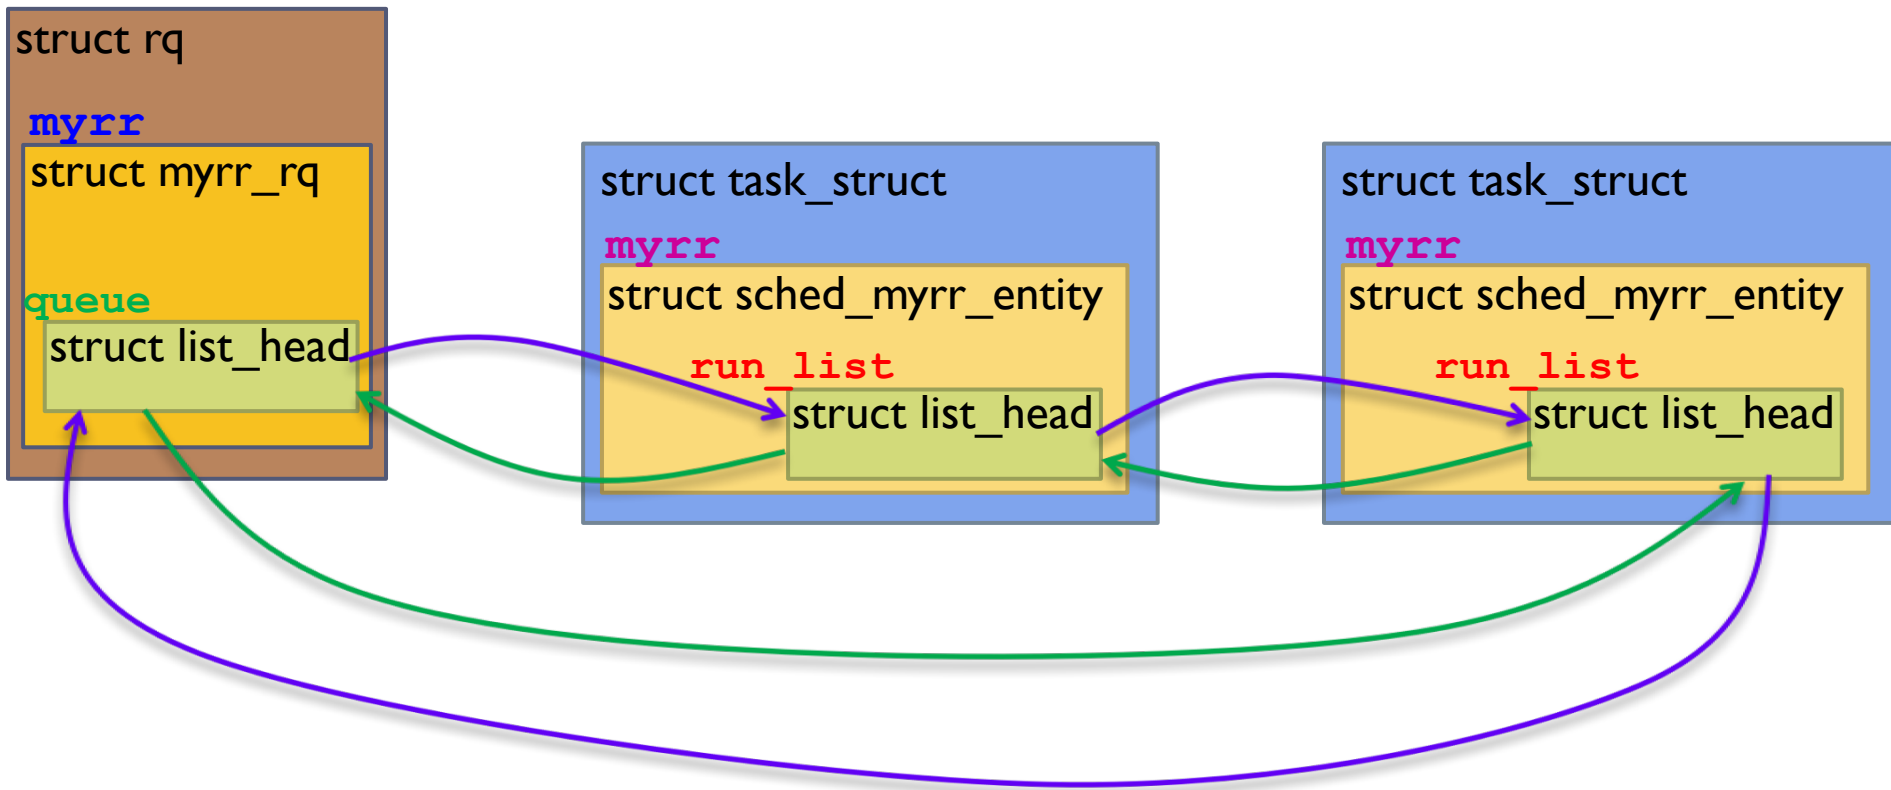
```
struct myrr_rq myrr;
```

after the line

```
struct rq_rq rt;
```

# Structure

---





## Step-4: Add implement schedule\_myrr.c

---

- ▶ Add the definition of myrr\_sched\_class in `kernel/sched/sched.h`
  - ▶ Add `extern const struct sched_class myrr_sched_class`
  - ▶ After `extern const struct sched_class idle_sched_class`

# Step-4: implement sched\_myrr.c

- Create a new file `kernel/sched/sched_myrr.c`

```
#include "sched.h"
static const struct sched_class myrr_sched_class = {
    .next = &fair_sched_class,
    .enqueue_task = enqueue_task_myrr,
    .dequeue_task = dequeue_task_myrr,
    .yield_task = yield_task_myrr,
    .check_preempt_curr = check_preempt_curr_myrr,
    .pick_next_task = pick_next_task_myrr,
    .put_prev_task = put_prev_task_myrr,
```

```
#ifdef CONFIG_SMP
    .select_task_rq = select_task_rq_myrr,
    .load_balance = load_balance_myrr,
    .move_one_task = move_one_task_myrr,
#endif
```

```
    .set_curr_task = set_curr_task_myrr,
    .task_tick = task_tick_myrr,
    .get_rr_interval = get_rr_interval_myrr,
    .prio_changed = prio_changed_myrr,
    .switched_to = switched_to_myrr,
};
```

We should have disabled SMP in the previous lab, we can ignore this part



## Step-4: implement sched\_myrr.c

---

```
static void enqueue_task_myrr(struct rq *rq, struct task_struct *p,
                             int wakeup, bool head)
{
    struct sched_myrr_entity *myrr_se = &p->myrr;
    list_add_tail(&myrr_se->run_list, &rq->myrr.queue);

    printk(KERN_INFO"[SCHED_MYRR] ENQUEUE: Process-%d\n", p->pid);
}

static void dequeue_task_myrr(struct rq *rq,
                              struct task_struct *p, int sleep)
{
    struct sched_myrr_entity *myrr_se = &p->myrr;
    list_del(&myrr_se->run_list);
    printk(KERN_INFO"[SCHED_MYRR] DEQUEUE: Process-%d\n", p->pid);
}
```



## Step-4: implement sched\_myrr.c

---

```
static void yield_task_myrr(struct rq *rq)
{
    struct sched_myrr_entity *myrr_se = &rq->curr->myrr;
    struct myrr_rq *myrr_rq = &rq->myrr;
    list_move_tail(&myrr_se->run_list, &myrr_rq->queue);
}

static void check_preempt_curr_myrr(struct rq *rq,
    struct task_struct *p, int flags)
{
    if (rq->curr->policy == SCHED_FIFO || rq->curr->policy == SCHED_RR)
        return ;
    if (rq->curr->policy == SCHED_MYRR)
        return ;
    /* preempt normal tasks */
    resched_task(rq->curr);
}
```

Basically the same as  
MY\_FIFO

## Step-4: implement sched\_myrr.c

---

```
static struct task_struct *pick_next_task_myrr(struct rq *rq)
{
    struct sched_myrr_entity *myrr_se = NULL;
    struct task_struct *p = NULL;
    struct myrr_rq *myrr_rq = &rq->myrr;
    if (list_empty(&myrr_rq->queue))
        return NULL;
    myrr_se = list_entry(myrr_rq->queue.next,
                        struct sched_myrr_entity,
                        run_list);
    p = container_of(myrr_se, struct task_struct, myrr);
    return p;
}

static void put_prev_task_myrr(struct rq *rq, struct task_struct *p)
{
    /* it is the place to update the current task's
     * runtime statistics */
}
```

## Step-4: implement sched\_myrr.c

---

```
#ifdef CONFIG_SMP
static int select_task_rq_myrr(struct rq *rq, struct task_struct *p,
                              int sd_flag, int flags)
{
    return task_cpu(p);
}

static unsigned long load_balance_myrr(struct rq *this_rq,
                                       int this_cpu, struct rq *busiest, unsigned long max_load_move,
                                       struct sched_domain *sd, enum cpu_idle_type idle,
                                       int *all_pinned, int *this_best_prio)
{
    return 0;
}

static int move_one_task_myrr(struct rq *this_rq, int this_cpu,
                              struct rq *busiest, struct sched_domain *sd,
                              enum cpu_idle_type idle)
{
    return 0;
}
#endif
```

We should have disabled SMP in the previous lab, we can ignore this part



## Step-4: implement sched\_myrr.c

```
static void set_curr_task_myrr(struct rq *rq)
{
}

static void task_tick_myrr(struct rq *rq,
                          struct task_struct *p, int queued)
{
    if (p->policy != SCHED_MYRR) return ;

    /* Do your implementation here!!!! */
    /* Hint: make use of time_slice attribute
       Hint: How can you access this attribute via p?
       Hint: How can you move a task to the end of the queue?
       Hint: Once the task is moved, a special function call
           set_tsk_need_resched(p) is required to be invoked!
    */
}
```

Function	Purpose
set_tsk_need_resched(p)	Set the need_resched flag in the given process

## Step-4: implement sched\_myrr.c

---

```
unsigned int get_rr_interval_myrr(struct rq *rq,  
                                struct task_struct *p)  
{  
    /* ToDo: return a default timeslice */  
}  
  
static void prio_changed_myrr(struct rq *rq, struct task_struct *p,  
                              int oldprio, int running)  
{  
}  
  
static void switched_to_myrr(struct rq *rq, struct task_struct *p,  
                              int running)  
{  
}
```



## Step-5: merge into the kernel

---

- ▶ Please follow the similar steps in the previous lab (i.e. `SCHED_MYFIFO`) to merge the new scheduler to the kernel
- ▶ Do we need to change the Makefile(s)?
- ▶ The steps may not be exactly the same (e.g. you need to change from `myfifo` to `myrr` in many places)
- ▶ Important:
  - ▶ When you initialize `myrr_entity` inside the `__sched_fork` function, you also need to initialize the attribute of `time_slice`
  - ▶ For example:
    - ▶ `INIT_LIST_HEAD(&p->myrr.run_list);`
    - ▶ `p->myrr.time_slice = /* Your default time slice */;`

## Step-6: write the test program

---

- ▶ Make some changes to the test program in the previous lab.

For example:

- ▶ `ret = sched_setscheduler(0, SCHED_MYFIFO, &sp);`
- ▶ `/* Now, we should replace it by what? */`

# Lab demo:

---

- ▶ After running the test program:
  - ▶ Show the kernel log message (we should have added appropriate `printk` function calls in `enqueue/dequeue` functions) to demonstrate that the round-robin scheduler is executed